CS 224n: Assignment #4

This assignment is split into two sections: *Neural Machine Translation with RNNs* and *Analyzing NMT Systems*. The first is primarily coding and implementation focused, whereas the second entirely consists of written, analysis questions. If you get stuck on the first section, you can always work on the second as the two sections are independent of each other. Note that the NMT system is more complicated than the neural networks we have previously constructed within this class and takes about **2 hours to train on a GPU**. Thus, we strongly recommend you get started early with this assignment. Finally, the notation and implementation of the NMT system is a bit tricky, so if you ever get stuck along the way, please come to Office Hours so that the TAs can support you.

1. Neural Machine Translation with RNNs (45 points)

In Machine Translation, our goal is to convert a sentence from the *source* language (e.g. Mandarin Chinese) to the *target* language (e.g. English). In this assignment, we will implement a sequence-to-sequence (Seq2Seq) network with attention, to build a Neural Machine Translation (NMT) system. In this section, we describe the **training procedure** for the proposed NMT system, which uses a Bidirectional LSTM Encoder and a Unidirectional LSTM Decoder.
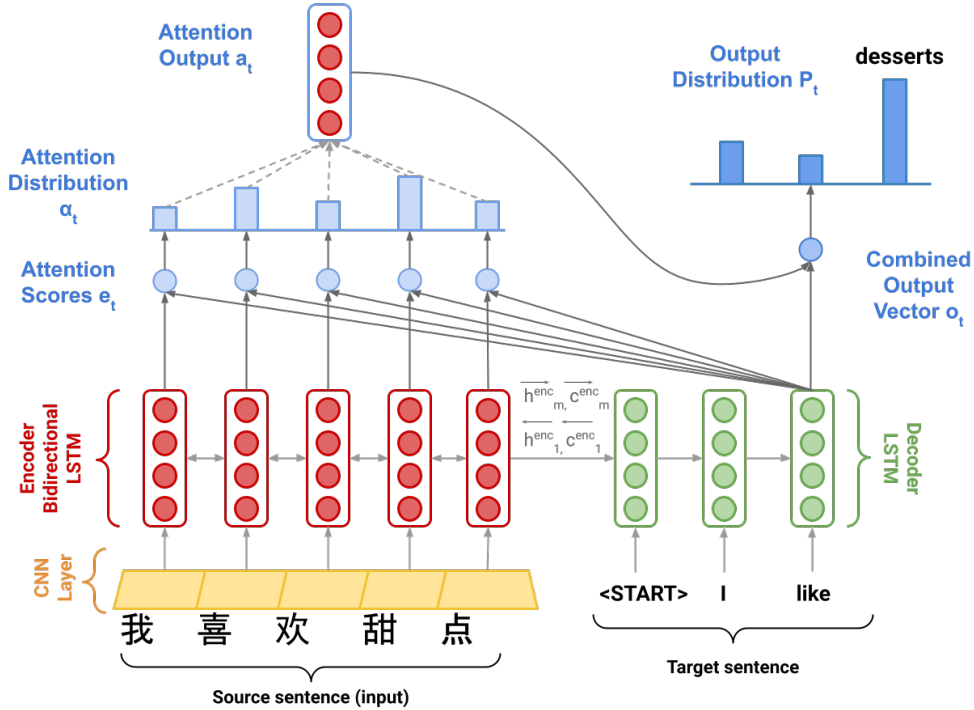

Figure 1: Seq2Seq Model with Multiplicative Attention, shown on the third step of the decoder. Hidden states $\mathbf{h}_i^{\text{enc}}$ and cell states $\mathbf{c}_i^{\text{enc}}$ are defined on the next page.

Model description (training procedure)

Given a sentence in the source language, we look up the character or word embeddings from an **embeddings matrix**, yielding $\mathbf{x}_1, \ldots, \mathbf{x}_m$ ($\mathbf{x}_i \in \mathbb{R}^{e \times 1}$), where m is the length of the source sentence and e is

the embedding size. We then feed the embeddings to a **convolutional layer** [1] while maintaining their shapes. We feed the convolutional layer outputs to the **bidirectional encoder**, yielding hidden states and cell states for both the forwards ($\rightarrow$) and backwards ($\leftarrow$) LSTMs. The forwards and backwards versions are concatenated to give hidden states $\mathbf{h}_i^{\text{enc}}$ and cell states $\mathbf{c}_i^{\text{enc}}$:

$$\mathbf{h}_i^{\text{enc}} = [\overleftarrow{\mathbf{h}_i^{\text{enc}}}; \overrightarrow{\mathbf{h}_i^{\text{enc}}}] \ \text{ where } \ \mathbf{h}_i^{\text{enc}} \in \mathbb{R}^{2h \times 1}, \overleftarrow{\mathbf{h}_i^{\text{enc}}}, \overrightarrow{\mathbf{h}_i^{\text{enc}}} \in \mathbb{R}^{h \times 1} \qquad 1 \leq i \leq m \qquad (1)$$

$$\mathbf{c}_i^{\text{enc}} = [\overleftarrow{\mathbf{c}_i^{\text{enc}}}; \overrightarrow{\mathbf{c}_i^{\text{enc}}}] \ \text{ where } \ \mathbf{c}_i^{\text{enc}} \in \mathbb{R}^{2h \times 1}, \overleftarrow{\mathbf{c}_i^{\text{enc}}}, \overrightarrow{\mathbf{c}_i^{\text{enc}}} \in \mathbb{R}^{h \times 1} \qquad 1 \leq i \leq m \qquad (2)$$

We then initialize the **decoder**'s first hidden state $\mathbf{h}_0^{\text{dec}}$ and cell state $\mathbf{c}_0^{\text{dec}}$ with a linear projection of the encoder's final hidden state and final cell state.[2]

$$\mathbf{h}_0^{\text{dec}} = \mathbf{W}_h[\overleftarrow{\mathbf{h}_1^{\text{enc}}}; \overrightarrow{\mathbf{h}_m^{\text{enc}}}] \ \text{ where } \ \mathbf{h}_0^{\text{dec}} \in \mathbb{R}^{h \times 1}, \mathbf{W}_h \in \mathbb{R}^{h \times 2h} \qquad (3)$$

$$\mathbf{c}_0^{\text{dec}} = \mathbf{W}_c[\overleftarrow{\mathbf{c}_1^{\text{enc}}}; \overrightarrow{\mathbf{c}_m^{\text{enc}}}] \ \text{ where } \ \mathbf{c}_0^{\text{dec}} \in \mathbb{R}^{h \times 1}, \mathbf{W}_c \in \mathbb{R}^{h \times 2h} \qquad (4)$$

With the decoder initialized, we must now feed it a target sentence. On the t^{th} step, we look up the embedding for the t^{th} subword, $\mathbf{y}_t \in \mathbb{R}^{e \times 1}$. We then concatenate $\mathbf{y}_t$ with the *combined-output vector* $\mathbf{o}_{t-1} \in \mathbb{R}^{h \times 1}$ from the previous timestep (we will explain what this is later down this page!) to produce $\overline{\mathbf{y}_t} \in \mathbb{R}^{(e+h) \times 1}$. Note that for the first target subword (i.e. the start token) $\mathbf{o}_0$ is a zero-vector. We then feed $\overline{\mathbf{y}_t}$ as input to the decoder.

$$\mathbf{h}_t^{\text{dec}}, \mathbf{c}_t^{\text{dec}} = \text{Decoder}(\overline{\mathbf{y}_t}, \mathbf{h}_{t-1}^{\text{dec}}, \mathbf{c}_{t-1}^{\text{dec}}) \ \text{ where } \ \mathbf{h}_t^{\text{dec}} \in \mathbb{R}^{h \times 1}, \mathbf{c}_t^{\text{dec}} \in \mathbb{R}^{h \times 1} \qquad (5)$$

$$(6)$$

We then use $\mathbf{h}_t^{\text{dec}}$ to compute multiplicative attention over $\mathbf{h}_1^{\text{enc}}, \ldots, \mathbf{h}_m^{\text{enc}}$:

$$\mathbf{e}_{t,i} = (\mathbf{h}_t^{\text{dec}})^T \mathbf{W}_{\text{attProj}} \mathbf{h}_i^{\text{enc}} \ \text{ where } \ \mathbf{e}_t \in \mathbb{R}^{m \times 1}, \mathbf{W}_{\text{attProj}} \in \mathbb{R}^{h \times 2h} \qquad 1 \leq i \leq m \qquad (7)$$

$$\alpha_t = \text{softmax}(\mathbf{e}_t) \ \text{ where } \ \alpha_t \in \mathbb{R}^{m \times 1} \qquad (8)$$

$$\mathbf{a}_t = \sum_{i=1}^{m} \alpha_{t,i} \mathbf{h}_i^{\text{enc}} \ \text{ where } \ \mathbf{a}_t \in \mathbb{R}^{2h \times 1} \qquad (9)$$

$\mathbf{e}_{t,i}$ is a scalar, the ith element of $\mathbf{e}_t \in \mathbb{R}^{m \times 1}$, computed using the hidden state of the decoder at the tth step, $\mathbf{h}_t^{\text{dec}} \in \mathbb{R}^{h \times 1}$, the attention projection $\mathbf{W}_{\text{attProj}} \in \mathbb{R}^{h \times 2h}$, and the hidden state of the encoder at the ith step, $\mathbf{h}_i^{\text{enc}} \in \mathbb{R}^{2h \times 1}$.

We now concatenate the attention output $\mathbf{a}_t$ with the decoder hidden state $\mathbf{h}_t^{\text{dec}}$ and pass this through a linear layer, tanh, and dropout to attain the *combined-output* vector $\mathbf{o}_t$.

$$\mathbf{u}_t = [\mathbf{a}_t; \mathbf{h}_t^{\text{dec}}] \ \text{ where } \ \mathbf{u}_t \in \mathbb{R}^{3h \times 1} \qquad (10)$$

$$\mathbf{v}_t = \mathbf{W}_u \mathbf{u}_t \ \text{ where } \ \mathbf{v}_t \in \mathbb{R}^{h \times 1}, \mathbf{W}_u \in \mathbb{R}^{h \times 3h} \qquad (11)$$

$$\mathbf{o}_t = \text{dropout}(\tanh(\mathbf{v}_t)) \ \text{ where } \ \mathbf{o}_t \in \mathbb{R}^{h \times 1} \qquad (12)$$

---

[1] Checkout https://cs231n.github.io/convolutional-networks for an in-depth description for convolutional layers if you are not familiar

[2] If it's not obvious, think about why we regard $[\overleftarrow{\mathbf{h}_1^{\text{enc}}}, \overrightarrow{\mathbf{h}_m^{\text{enc}}}]$ as the 'final hidden state' of the Encoder.

Then, we produce a probability distribution $\mathbf{P}_t$ over target subwords at the t^{th} timestep:

$$\mathbf{P}_t = \text{softmax}(\mathbf{W}_{\text{vocab}}\mathbf{o}_t) \ \text{ where } \ \mathbf{P}_t \in \mathbb{R}^{V_t \times 1}, \mathbf{W}_{\text{vocab}} \in \mathbb{R}^{V_t \times h} \tag{13}$$

Here, V_t is the size of the target vocabulary. Finally, to train the network we then compute the cross entropy loss between $\mathbf{P}_t$ and $\mathbf{g}_t$, where $\mathbf{g}_t$ is the one-hot vector of the target subword at timestep t:

$$J_t(\theta) = \text{CrossEntropy}(\mathbf{P}_t, \mathbf{g}_t) \tag{14}$$

Here, θ represents all the parameters of the model and $J_t(\theta)$ is the loss on step t of the decoder. Now that we have described the model, let's try implementing it for Mandarin Chinese to English translation!

Setting up your Virtual Machine

Follow the instructions in the CS224n Azure Guide (link also provided on website and Ed) in order to create your VM instance. This should take you approximately 45 minutes. Though you will need the GPU to train your model, we strongly advise that you first develop the code locally and ensure that it runs, before attempting to train it on your VM. GPU time is expensive and limited. It takes approximately **1.5 to 2 hours** to train the NMT system. We don't want you to accidentally use all your GPU time for debugging your model rather than training and evaluating it. Finally, **make sure that your VM is turned off whenever you are not using it.**

If your Azure subscription runs out of money, your VM will be temporarily locked and inaccessible. If that happens, please fill out a request form here.

In order to run the model code on your **local** machine, please run the following command to create the proper virtual environment:

```
conda env create --file local_env.yml
```

Note that this virtual environment **will not** be needed on the VM.

Implementation and written questions

(a) (2 points) (coding) In order to apply tensor operations, we must ensure that the sentences in a given batch are of the same length. Thus, we must identify the longest sentence in a batch and pad others to be the same length. Implement the pad_sents function in utils.py, which shall produce these padded sentences.

(b) (3 points) (coding) Implement the __init__ function in model_embeddings.py to initialize the necessary source and target embeddings.

(c) (4 points) (coding) Implement the __init__ function in nmt_model.py to initialize the necessary model layers (LSTM, CNN, projection, and dropout) for the NMT system.

(d) (8 points) (coding) Implement the encode function in nmt_model.py. This function converts the padded source sentences into the tensor $\mathbf{X}$, generates $\mathbf{h}_1^{\text{enc}}, \ldots, \mathbf{h}_m^{\text{enc}}$, and computes the initial state $\mathbf{h}_0^{\text{dec}}$ and initial cell $\mathbf{c}_0^{\text{dec}}$ for the Decoder. You can run a non-comprehensive sanity check by executing:

```
python sanity_check.py 1d
```

(e) (8 points) (coding) Implement the decode function in nmt_model.py. This function constructs $\bar{\mathbf{y}}$ and runs the step function over every timestep for the input. You can run a non-comprehensive sanity check by executing:

```
python sanity_check.py 1e
```

(f) (10 points) (coding) Implement the step function in nmt_model.py. This function applies the Decoder's LSTM cell for a single timestep, computing the encoding of the target subword $\mathbf{h}_t^{\text{dec}}$, the attention scores $\mathbf{e}_t$, attention distribution α_t, the attention output $\mathbf{a}_t$, and finally the combined output $\mathbf{o}_t$. You can run a non-comprehensive sanity check by executing:

```
python sanity_check.py 1f
```

(g) (3 points) (written) The generate_sent_masks() function in nmt_model.py produces a tensor called enc_masks. It has shape (batch size, max source sentence length) and contains 1s in positions corresponding to 'pad' tokens in the input, and 0s for non-pad tokens. Look at how the masks are used during the attention computation in the step() function (lines 311-312).

First explain (in around three sentences) what effect the masks have on the entire attention computation. Then explain (in one or two sentences) why it is necessary to use the masks in this way.

Solution: 1. **What effect:** for every batch, those attention scores which have corresponding zero-padded embeddings are set to $-\infty$. This way, during the calculation of *attention distributions* α_t, the probabilities are calculated over scores with corresponding non-padded words whereas the scores with corresponding padded words are negligible. Finally, during the calculation of *attention outputs A_t*, for every batch only the addition of those hidden states matter which are not multiplied by a probability close to 0.

2. **Why necessary:** using masks in this way is an efficient way to determine the true *attention distribution* that only involves the non-padded entries. Involving padded entries would result in false *attention* representation.

Now it's time to get things running! As noted earlier, we recommend that you develop the code on your personal computer. Confirm that you are running in the proper conda environment and then execute the following command to train the model on your local machine:

```
sh run.sh train_local
(Windows) run.bat train_local
```

For a faster way to debug by training on less data, you can run the following instead:

```
sh run.sh train_debug
(Windows) run.bat debug
```

To help with monitoring and debugging, the starter code uses tensorboard to log loss and perplexity during training using TensorBoard[3]. TensorBoard provides tools for logging and visualizing training information from experiments. To open TensorBoard, run the following in your conda environment:

```
tensorboard --logdir=runs
```

You should see a significant decrease in loss during the initial iterations. Once you have ensured that your code does not crash (i.e. let it run till iter 10 or iter 20), power on your VM from the Azure Web

---

[3]https://pytorch.org/docs/stable/tensorboard.html

Portal. Then read the *Managing Code Deployment to a VM* section of our Practical Guide to VMs (link also given on website and Ed) for instructions on how to upload your code to the VM.

Next, install necessary packages to your VM by running:

```
pip install -r gpu_requirements.txt
```

Finally, turn to the *Managing Processes on a VM* section of the Practical Guide and follow the instructions to create a new tmux session. Concretely, run the following command to create tmux session called nmt.

```
tmux new -s nmt
```

Once your VM is configured and you are in a tmux session, execute:

```
sh run.sh train
(Windows) run.bat train
```

Once you know your code is running properly, you can detach from session and close your ssh connection to the server. To detach from the session, run:

```
tmux detach
```

You can return to your training model by ssh-ing back into the server and attaching to the tmux session by running:

```
tmux a -t nmt
```

(h) (3 points) (written) Once your model is done training (**this should take under 2 hours on the VM**), execute the following command to test the model:

```
sh run.sh test
(Windows) run.bat test
```

Please report the model's corpus BLEU Score. It should be larger than 18.

> **Solution:**   Corpus BLEU: 21.014487279150515

(i) (4 points) (written) In class, we learned about dot product attention, multiplicative attention, and additive attention. As a reminder, dot product attention is $\mathbf{e}_{t,i} = \mathbf{s}_t^T \mathbf{h}_i$, multiplicative attention is $\mathbf{e}_{t,i} = \mathbf{s}_t^T \mathbf{W} \mathbf{h}_i$, and additive attention is $\mathbf{e}_{t,i} = \mathbf{v}^T \tanh(\mathbf{W}_1 \mathbf{h}_i + \mathbf{W}_2 \mathbf{s}_t)$.

  i. (2 points) Explain one advantage and one disadvantage of *dot product attention* compared to multiplicative attention.

  ii. (2 points) Explain one advantage and one disadvantage of *additive attention* compared to multiplicative attention.

> **Solution:**   i.) An advantage of dot product attention compared to multiplicative attention is that it does not have any learnable parameters and it is a vector dot product, and thus is very fast to compute. A simple dot product does not provide much information and is not sufficient to capture the intricacies on the other hand.
> ii.) The advantage is that the similarity is computed in a non-linearly transformed space. This adds more flexibility in terms of parameter space. The disadvantage is that computation is expensive.

2. Analyzing NMT Systems (25 points)

(a) (3 points) Look at the `src.vocab` file for some examples of phrases and words in the source language vocabulary. When encoding an input Mandarin Chinese sequence into "pieces" in the vocabulary, the tokenizer maps the sequence to a series of vocabulary items, each consisting of one or more characters (thanks to the `sentencepiece` tokenizer, we can perform this segmentation even when the original text has no white space). Given this information, how could adding a 1D Convolutional layer after the embedding layer and before passing the embeddings into the bidirectional encoder help our NMT system? **Hint:** each Mandarin Chinese character is either an entire word or a morpheme in a word. Look up the meanings of 电, 脑, and 电脑 separately for an example. The characters 电 (electricity) and 脑 (brain) when combined into the phrase 电脑 mean computer.

Solution: In Neural Machine translation (NMT), the role of the tokenizer is to break down the input sequence into smaller pieces or tokens. Adding a 1D Concolutional layer after the embedding can help the NMT system capture local dependencies within the input sequence. The layer acts like a feature extractor tat identifies and preserves these local decencies before the sequence is passed in to the bidirectional encoder.

(b) (8 points) Here we present a series of errors we found in the outputs of our NMT model (which is the same as the one you just trained). For each example of a reference (i.e., 'gold') English translation, and NMT (i.e., 'model') English translation, please:

1. Identify the error in the NMT translation.
2. Provide possible reason(s) why the model may have made the error (either due to a specific linguistic construct or a specific model limitation).
3. Describe one possible way we might alter the NMT system to fix the observed error. There are more than one possible fixes for an error. For example, it could be tweaking the size of the hidden layers or changing the attention mechanism.

Below are the translations that you should analyze as described above. Only analyze the underlined error in each sentence. Rest assured that you don't need to know Mandarin to answer these questions. You just need to know English! If, however, you would like some additional color on the source sentences, feel free to use a resource like https://www.archchinese.com/chinese_english_dictionary.html to look up words. Feel free to search the training data file to have a better sense of how often certain characters occur.

i. (2 points) **Source Sentence:** 贼人其后被警方拘捕及被判处盗窃罪名成立。
Reference Translation: *the culprits were subsequently arrested and convicted.*
NMT Translation: *the culprit was subsequently arrested and sentenced to theft.*

Solution: 1.
Error Identification: The NMT translation incorrectly translates "贼人" as "the culprit" (singular) instead of "the culprits" (plural) as in the reference translation.
2. **Possible Reason:** This could be due to the fact that in Mandarin, plurality is not always explicitly marked on the noun. The model might have learned to default to the singular form when the plurality is not explicitly marked in the source sentence.
3. **Possible Fix:** One possible way to fix this error could be to incorporate a better handling of plurality in the model. This could be achieved by using a more sophisticated attention mechanism that can better capture such nuances in the source sentence. Additionally, using a larger training dataset that includes more examples of plural nouns could also help the model learn this aspect of the language better.

ii. (2 points) **Source Sentence**: 几乎已经没有地方容纳这些人, 资源已经用尽。
Reference Translation: *there is almost no space to accommodate these people, and resources have run out.*
NMT Translation: *the resources have been exhausted and <u>resources have been exhausted</u>.*

Solution:
1. **Error Identification:** The NMT translation incorrectly repeats the phrase "resources have been exhausted" and completely misses the first part of the sentence about there being almost no space to accommodate these people.
2. **Possible Reason:** This could be due to an error in the attention mechanism of the model where it incorrectly gave more weight to the latter part of the sentence about resources being exhausted and missed the first part of the sentence.
3. **Possible Fix:** One possible way to fix this error could be to improve the attention mechanism of the model so that it can better capture all parts of the source sentence. Additionally, using a larger and more diverse training dataset could also help the model learn to handle such sentences better. Another possible solution could be to incorporate a mechanism in the model that penalizes repeated phrases to avoid such repetition errors.

iii. (2 points) **Source Sentence**: 当局已经宣布今天是国殇日。
Reference Translation: *authorities have announced <u>a national mourning today.</u>*
NMT Translation: *the administration has announced <u>today's day.</u>*

Solution:
1. **Error Identification:** The NMT translation incorrectly translates " 国殇日" as "today's day" instead of "a national mourning today" as in the reference translation.
2. **Possible Reason:** This could be due to the model not correctly learning the mapping of " 国殇日" to "a national mourning day". It might be because such phrases are less common and thus the model has not seen enough examples during training to learn this correctly.
3. **Possible Fix:** One possible way to fix this error could be to include more examples of such phrases in the training data so that the model can learn the correct mapping. Additionally, using a larger and more diverse training dataset could also help the model learn to handle such sentences better. Another possible solution could be to incorporate a mechanism in the model that penalizes incorrect translations to avoid such errors.

iv. (2 points) **Source Sentence[4]:** 俗语有云:"唔做唔错"。
Reference Translation: *<u>" act not, err not ",</u> so a saying goes.*
NMT Translation: *as the saying goes, <u>" it's not wrong. "</u>*

Solution:
1. **Error Identification:** The NMT translation incorrectly translates " 唔做唔错" as "it's not wrong" instead of "act not, err not" as in the reference translation.
2. **Possible Reason:** This could be due to the model not correctly learning the mapping of " 唔做唔错" to "act not, err not". It might be because such phrases are less common and thus the model has not seen enough examples during training to learn this correctly.
3. **Possible Fix:** One possible way to fix this error could be to include more examples of such phrases in the training data so that the model can learn the correct mapping. Additionally,

---

[4]This is a Cantonese sentence! The data used in this assignment comes from GALE Phase 3, which is a compilation of news written in simplified Chinese from various sources scraped from the internet along with their translations. For more details, see https://catalog.ldc.upenn.edu/LDC2017T02.

using a larger and more diverse training dataset could also help the model learn to handle such sentences better. Another possible solution could be to incorporate a mechanism in the model that penalizes incorrect translations to avoid such errors.

(c) (14 points) BLEU score is the most commonly used automatic evaluation metric for NMT systems. It is usually calculated across the entire test set, but here we will consider BLEU defined for a single example.[5] Suppose we have a source sentence $\mathbf{s}$, a set of k reference translations $\mathbf{r}_1, \ldots, \mathbf{r}_k$, and a candidate translation $\mathbf{c}$. To compute the BLEU score of $\mathbf{c}$, we first compute the *modified n-gram precision* p_n of $\mathbf{c}$, for each of $n = 1, 2, 3, 4$, where n is the n in n-gram:

$$p_n = \frac{\sum_{\text{ngram} \in \mathbf{c}} \min\left(\max_{i=1,\ldots,k} \text{Count}_{\mathbf{r}_i}(\text{ngram}), \ \text{Count}_{\mathbf{c}}(\text{ngram}) \right)}{\sum_{\text{ngram} \in \mathbf{c}} \text{Count}_{\mathbf{c}}(\text{ngram})} \tag{15}$$

Here, for each of the n-grams that appear in the candidate translation $\mathbf{c}$, we count the maximum number of times it appears in any one reference translation, capped by the number of times it appears in $\mathbf{c}$ (this is the numerator). We divide this by the number of n-grams in $\mathbf{c}$ (denominator).

Next, we compute the *brevity penalty* BP. Let $len(c)$ be the length of $\mathbf{c}$ and let $len(r)$ be the length of the reference translation that is closest to $len(c)$ (in the case of two equally-close reference translation lengths, choose $len(r)$ as the shorter one).

$$BP = \begin{cases} 1 & \text{if } len(c) \geq len(r) \\ \exp\left(1 - \frac{len(r)}{len(c)}\right) & \text{otherwise} \end{cases} \tag{16}$$

Lastly, the BLEU score for candidate $\mathbf{c}$ with respect to $\mathbf{r}_1, \ldots, \mathbf{r}_k$ is:

$$BLEU = BP \times \exp\left(\sum_{n=1}^{4} \lambda_n \log p_n \right) \tag{17}$$

where $\lambda_1, \lambda_2, \lambda_3, \lambda_4$ are weights that sum to 1. The log here is natural log.

  i. (5 points) Please consider this example:
  Source Sentence $\mathbf{s}$: **需要有充足和可预测的资源。**
  Reference Translation $\mathbf{r}_1$: *resources have to be sufficient and they have to be predictable*
  Reference Translation $\mathbf{r}_2$: *adequate and predictable resources are required*
  NMT Translation $\mathbf{c}_1$: there is a need for adequate and predictable resources
  NMT Translation $\mathbf{c}_2$: resources be sufficient and predictable to
  Please compute the BLEU scores for $\mathbf{c}_1$ and $\mathbf{c}_2$. Let $\lambda_i = 0.5$ for $i \in \{1,2\}$ and $\lambda_i = 0$ for $i \in \{3,4\}$ (**this means we ignore 3-grams and 4-grams**, i.e., don't compute p_3 or p_4). When computing BLEU scores, show your work (i.e., show your computed values for p_1, p_2, $len(c)$, $len(r)$ and BP). Note that the BLEU scores can be expressed between 0 and 1 or between 0 and 100. The code is using the 0 to 100 scale while in this question we are using the **0 to 1** scale. Please round your responses to 3 decimal places.

  Which of the two NMT translations is considered the better translation according to the BLEU Score? Do you agree that it is the better translation?

---

[5]This definition of sentence-level BLEU score matches the sentence_bleu() function in the nltk Python package. Note that the NLTK function is sensitive to capitalization. In this question, all text is lowercased, so capitalization is irrelevant. http://www.nltk.org/api/nltk.translate.html#nltk.translate.bleu_score.sentence_bleu

Solution: Here is a list of different n-grams (only considering $1 - 2$-grams) from NMT translations and their counts in brackets (in $\mathbf{c}$| in $\mathbf{r_1}$ | in $\mathbf{r_2}$):

1. $\mathbf{c_1}$
   - **1-grams**: *and* $(2|1|2)$, *the* $(3|3|3)$, *light* $(1|1|1)$, *shines* $(1|1|1)$, *in* $(1|1|1)$, *darkness* $(2|2|2)$, *can* $(1|0|0)$, *not* $(1|1|1)$, *comprehend* $(1|0|1)$.
   - **2-grams**: *and the* $(2|1|2)$, *the light* $(1|1|1)$, *light shines* $(1|1|1)$, *shines in* $(1|1|1)$, *in the* $(1|1|1)$, *the darkness* $(2|2|2)$, *darkness and* $(1|1|1)$, *darkness can* $(1|0|0)$, *can not* $(1|0|0)$, *not comprehend* $(1|0|1)$.

2. $\mathbf{c_2}$
   - **1-grams**: *the* $(4|3|3)$, *light* $(1|1|1)$, *shines* $(1|1|1)$, *darkness* $(2|2|2)$, *has* $(1|1|0)$, *not* $(1|1|1)$, *in* $(1|1|1)$, *and* $(1|1|2)$, *trials* $(1|0|0)$.
   - **2-grams**: *the light* $(1|1|1)$, *light shines* $(1|1|1)$, *shines the* $(1|0|0)$, *the darkness* $(2|2|2)$, *darkness has* $(1|1|0)$, *has not* $(1|1|0)$, *not in* $(1|0|0)$, *in the* $(1|1|1)$, *darkness and* $(1|1|1)$, *and the* $(1|1|2)$, *the trials* $(1|0|0)$.

p_1 and p_2 for both $\mathbf{c_1}$ and $\mathbf{c_2}$ are now easy to compute:

$$p_{1,\mathbf{c_1}} = \frac{2 + 3 + 1 + 1 + 1 + 2 + 0 + 1 + 1}{2 + 3 + 1 + 1 + 1 + 2 + 1 + 1 + 1} = \frac{12}{13} \approx 0.9231 \tag{18}$$

$$p_{2,\mathbf{c_1}} = \frac{2 + 1 + 1 + 1 + 1 + 2 + 1 + 0 + 0 + 1}{2 + 1 + 1 + 1 + 1 + 2 + 1 + 1 + 1 + 1} = \frac{10}{12} \approx 0.8333 \tag{19}$$

$$p_{1,\mathbf{c_2}} = \frac{3 + 1 + 1 + 2 + 1 + 1 + 1 + 1 + 0}{4 + 1 + 1 + 2 + 1 + 1 + 1 + 1 + 1} = \frac{11}{13} \approx 0.8462 \tag{20}$$

$$p_{2,\mathbf{c_2}} = \frac{1 + 1 + 0 + 2 + 1 + 1 + 0 + 1 + 1 + 1 + 0}{1 + 1 + 1 + 2 + 1 + 1 + 1 + 1 + 1 + 1 + 1} = \frac{9}{12} = 0.7500 \tag{21}$$

Now we can compute $len(c)$ and $len(r)$ for both $\mathbf{c_1}$ and $\mathbf{c_2}$. Note that in both cases $\mathbf{r_1}$ is closer to NMT translation lengths thus its length will be checked:

$$len(c)_{\mathbf{c_1}} = 13; \quad len(r)_{\mathbf{c_1}} = 13 \tag{22}$$

$$len(c)_{\mathbf{c_2}} = 13; \quad len(r)_{\mathbf{c_2}} = 13 \tag{23}$$

Since in both cases $len(c) = len(r)$, BP will be 1 for both $\mathbf{c_1}$ and $\mathbf{c_2}$:

$$BP_{\mathbf{c_1}} = 1 \tag{24}$$

$$BP_{\mathbf{c_2}} = 1 \tag{25}$$

Now we just substitute the numbers and compute BLEU score, given that $\lambda_1 = 0.5$ and $\lambda_2 = 0.5$ for both $\mathbf{c_1}$ and $\mathbf{c_2}$. Since $\lambda_3 = 0$ and $\lambda_4 = 0$ are for both $\mathbf{c_1}$ and $\mathbf{c_2}$, we just ignore those terms during summation operation:

$$BLEU_{\mathbf{c_1}} = BP_{\mathbf{c_1}} \times \exp\left(\lambda_1 \log p_{1,\mathbf{c_1}} + \lambda_2 \log p_{2,\mathbf{c_1}}\right) \approx 0.8771 \tag{26}$$

$$BLEU_{\mathbf{c_2}} = BP_{\mathbf{c_2}} \times \exp\left(\lambda_1 \log p_{1,\mathbf{c_2}} + \lambda_2 \log p_{2,\mathbf{c_2}}\right) \approx 0.7966 \tag{27}$$

The better translation is $\mathbf{c_1}$ since $BLEU_{\mathbf{c_1}} > BLEU_{\mathbf{c_2}}$. This is indeed reflected in the translation as it is more similar to the reference sentences and has a more logical flow than $\mathbf{c_2}$.

ii. (5 points) Our hard drive was corrupted and we lost Reference Translation $\mathbf{r_1}$. Please recompute BLEU scores for $\mathbf{c_1}$ and $\mathbf{c_2}$, this time with respect to $\mathbf{r_2}$ only. Which of the two NMT translations now receives the higher BLEU score? Do you agree that it is the better translation?

Solution: Keeping the same lists of n-grams, it is easy to recompute p_1 and p_2 for $\mathbf{c}_1$ and $\mathbf{c}_2$:

$$p_{1,\mathbf{c}_1} = \frac{1+3+1+1+1+2+0+1+0}{2+3+1+1+1+2+1+1+1} = \frac{10}{13} \approx 0.7692 \tag{28}$$

$$p_{2,\mathbf{c}_1} = \frac{1+1+1+1+1+2+1+0+0+0}{2+1+1+1+1+2+1+1+1+1} = \frac{8}{12} \approx 0.6667 \tag{29}$$

$$p_{1,\mathbf{c}_2} = \frac{3+1+1+2+1+1+1+1+0}{4+1+1+2+1+1+1+1+1} = \frac{11}{13} \approx 0.8462 \tag{30}$$

$$p_{2,\mathbf{c}_2} = \frac{1+1+0+2+1+1+0+1+1+1+0}{1+1+1+2+1+1+1+1+1+1+1} = \frac{9}{12} \approx 0.7500 \tag{31}$$

$len(c)$ and $len(r)$ for both $\mathbf{c}_1$ and $\mathbf{c}_2$ will be the same because there is only $\mathbf{r}_1$ which is the same as when both sentences were available. This means that both BC are also the same.

$$BP_{\mathbf{c}_1} = 1 \tag{32}$$

$$BP_{\mathbf{c}_2} = 1 \tag{33}$$

Finally, we can recompute the BLEU score by applying the same formula and the same λ constants:

$$BLEU_{\mathbf{c}_1} = BP_{\mathbf{c}_1} \times \exp\left(\lambda_1 \log p_{1,\mathbf{c}_1} + \lambda_2 \log p_{2,\mathbf{c}_1}\right) \approx 0.7161 \tag{34}$$

$$BLEU_{\mathbf{c}_2} = BP_{\mathbf{c}_2} \times \exp\left(\lambda_1 \log p_{1,\mathbf{c}_2} + \lambda_2 \log p_{2,\mathbf{c}_2}\right) \approx 0.7966 \tag{35}$$

Now the second translation $\mathbf{c}_1$ receives the higher score. That is because of more n-gram co-ocurances between $\mathbf{c}_1$ and $\mathbf{r}_1$ than between $\mathbf{c}_2$ and $\mathbf{r}_1$. However, those co-ocurances are not in the similar places - in $\mathbf{c}_2$ they do not seem to have a logical structure (e.g., "the darkness has not in the darkness"). This could be solved by taking into account 3-grams and 4-grams. Thus the better score does not justify the translation.

iii. (2 points) Due to data availability, NMT systems are often evaluated with respect to only a single reference translation. Please explain (in a few sentences) why this may be problematic. In your explanation, discuss how the BLEU score metric assesses the quality of NMT translations when there are multiple reference transitions versus a single reference translation.

Solution: When there are multiple reference translations, maximum number of some particular n-grams is taken across all references when computing the summation term in the numerator of p_n meaning the captured n-gram in the NMT translation is given more chances to occur at least the same number of times as in the NMT translation. That way p_n is not reduced and the higher it is, the higher BLEU is. This suggests that NMT translation could be very valid, it just may be some form of alternative and not considering that may lead to not very accurate BLEU scores.

iv. (2 points) List two advantages and two disadvantages of BLEU, compared to human evaluation, as an evaluation metric for Machine Translation.

Solution:
1. **Advantages**:
   (a) *Automated* - does not need human resources, is reliable and does not make errors
   (b) *Fast* - evaluation for a single NMT translation is very fast
   (c) *Language-independent* - same metrics works for every language

    2. **Disadvantages**:

     (a) *Structure-independent* - it does not consider the natural flow of the sentences, the score is mainly based on n-gram co-ocurrance counts which could happen randomly in sentences

     (b) *Resource-dependent* - requires at least several reference sentences in order to properly evaluate the NMT translation

Submission Instructions

You shall submit this assignment on GradeScope as two submissions – one for "Assignment 4 [coding]" and another for 'Assignment 4 [written]":

1. Run the collect_submission.sh script on Azure to produce your assignment4.zip file. You can use scp to transfer files between Azure and your local computer.

2. Upload your assignment4.zip file to GradeScope to "Assignment 4 [coding]".

3. Upload your written solutions to GradeScope to "Assignment 4 [written]". When you submit your assignment, make sure to tag all the pages for each problem according to Gradescope's submission directions. Points will be deducted if the submission is not correctly tagged.